

Computing Science (CMPUT) 455

Search, Knowledge, and Simulations

Ting-Han Wei

Department of Computing Science
University of Alberta
`tinghan@ualberta.ca`

Fall 2020

Lecture 4: Formalizing Decision-Making

Lecture Topics

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Intro to computer decision-making
- Search, state spaces, state space of a game
- Terminology for state space search
- Types and models of state spaces
- Game trees, b^d model

Decision-making Exercise - Use Computers?

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Could a computer do the same decision-making processes as you?
- What would a computer need to make similar decisions?
 - Knowledge about your problem?
 - Knowledge about the world, “common sense”?
 - Logical reasoning?
 - Optimization?
 - Input? Sensors? Memory? Processing power?

Do you Want Computers to Make Decisions?

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Would you want a computer to make decisions **for you**?
- When would you trust it?
- How about computers supporting human decision-making?
- Who is in control?
 - Machine?
 - Programmer?
 - Employer?
 - Government?
 - Nobody?

Reasoning - Now vs Future

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Some decision-making considers only the present
- Reflexes
- Intuitive Decisions
- Tasks where looking ahead is not needed
- Example: image recognition
 - All relevant knowledge is available now
 - Neural net has learned knowledge from many (past) examples

Reasoning - Now vs Future

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Much reasoning involves thinking about the future, in order to make decisions now
- Example: should I take an umbrella to work?
 - It does not rain now
 - The forecast predicts rain later today
- Chess example: should I capture this queen?
 - It looks good now
 - It gets me checkmated 8 moves later

Decision Making as Search Problem

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Decision making in practice has many complications
- Often, we do not know:
 - The possible actions
 - The goals of others
 - How to evaluate different outcomes
 - How to compute with limited resources
 - ...

*It is difficult to make predictions,
especially about the future.*

<https://quoteinvestigator.com/2013/10/20/no-predict/>

- Where do we start? Simplify. A lot.

A Simple Setting for Decision-Making

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

Long list of assumptions:

- The *state* of the world is completely known at each time
- There are *terminal states* where we can evaluate the result precisely, e.g.
 - A number: value, score, reward, utility,...
 - One of a set of possible outcomes, e.g. {win, loss, draw}
- The possible *actions* (or moves) in each state are known
- An action changes the state to a new state in a known, deterministic way
- No chance element -
no dice rolls, cards drawn, other random events

Do games fit this simple setting?

Do Games fit Our Simple Setting?

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Single-player (puzzles): often, yes
 - Rubik's cube, solitaire, rushhour, sudoku, crossword puzzles, 15-puzzle
- **Two-player games**: our main topic - see separate slide
- Multi-player games: mixed
 - Yes: Chinese checkers, ...
 - No: most card games,
all dice games,
most family board games, ...

Do Games fit Our Simple Setting?

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Online and computer games: mostly, no
 - Simultaneous moves, often in real time
 - Map only partially known
 - Complex physics simulations
 - Many other complications...

“Classical” Two Player Games

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Our simple setting, plus:
- Two players, often called Black and White
- Players move alternately: I play, you play, I play,...
- A move instantly changes the state (no duration, no slow transitions)
- Simplest, most frequent case is *zero-sum*:
my win is opponent's loss
 - Opposite: cooperative games
- Examples: chess, checkers, Go, Tic-Tac-Toe, ...

Why Study Decision-Making using “Classical” Games?

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Simple, controlled environment
- Still hard to solve or play well
- Interesting for many people
- Games and results are easy to understand
- Playing games well requires good decision-making skills
- We can study the *core problems* of decision-making without being distracted by too many complications

Formalizing State Space Search

Terminology

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

The next few slides introduce the terms we use to talk about state space search in general, and specifically, games.

- game state, state
- state space, game graph, game tree
- board state, position
- move, action
- move sequence, history, game record
- score, value, evaluation, result

State, Game State

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Complete description of the current situation
- In games: board position or cards etc, toPlay (whose turn it is)
- State, plus rules of game, allow us to determine actions (moves)
- Often includes (parts of) history:
 - Sequence of moves from start of game to current position

Example in Go1

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- How is game state represented in Go1?
- In `simple_board.py`
- Class `SimpleGoBoard`
- Contains 1-d array `board`
 - Each array entry contains the color of one point
- Contains field `ko_recapture` to implement simple ko rule
- `current_player` (BLACK or WHITE)
- Other fields: `name`, `version`,...

Action, Move

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Leads from one state to another
- Move in games may include `toPlay`:
color of the player who plays the move
- *Alternating* play:
 - `current_player` changes after each move
 - We do not need to specify color with the move (e.g. can just store it in the state)
- Move sequence, history, game record:
all moves played in a game
- In `Go1`: move represented by index of point in array,
and color
 - Used e.g. in `play_move(point, color)`

State vs History

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- When is history needed?
 - Depends on rules, structure of search space
- Example: Ko (repetition), legal moves in Go depend on previous move history
- Example: TicTacToe does *not* need history
- Compare with *Markov Decision Problems* (MDP) in single agent search, *Markov property*:
 - History is irrelevant in MDP
 - Current state contains all relevant information

History-only State?

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- How about having *only* history, no other information in state?
- Yes, that works in principle
 - Rules plus complete history determine the state exactly
 - Examples:
 - Game records with list of moves
 - Sequence of GTP play commands
- It may be inefficient if we always have to replay all actions from the beginning in order to find the current state
- Some forms of machine learning work with this representation

State Space

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- A *state space* is:
 - A graph with all the possible states of a problem
 - Edges in graph show how states are connected by actions
- State space represented as directed graph $G = (V, E)$:
- Nodes in V : game states
- Directed edges in E : moves
 - Edge $e = (s_1, s_2)$ contains:
 - State s_1 before move
 - State s_2 after move


State Space Representation - Go1 Example

CMPUT 455


Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces



board1



board2

- **State before move:**
`(board1, to_play = WHITE,
ko_recapture = None, ...)`
- **Play move: White B1**
- **State after move:**
`(board2, to_play = BLACK,
ko_recapture = None, ...)`

Terminal State

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- A *terminal state* has no possible moves (actions)
- No outgoing edges in graph
- The rules of a game decide:
- When is the game over?
(did we reach a terminal state?)
- What is the outcome in a terminal state?

Terminal State in Go

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Game can end in one of two ways
 - A player resigns
 - Both players pass in turn
- Most Go players do not keep playing until there are no legal moves
 - Some moves are bad anyway and should not be played (see `Go0` vs `Go1`)
- Stop playing when the ownership of each point is clear to both players, then count the score
- Example: `Go1.py` stops playing when there are only single point eyes left

Terminal States and Rewards

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Later, in reinforcement learning, we will talk a lot about *rewards* (or costs: negative rewards)
- In many games, the only reward is at the end, in the terminal state
- Example: +1 if you win, -1 if you lose
- A few games have other, earlier rewards/costs
 - Example: blinds and bidding in poker

Types of State Spaces

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces




Image source:

sciencefair.math.iit.edu

- Assume *root* at the top is current state

- 1 Tree
- 2 DAG (directed acyclic graph)
- 3 DCG (directed cyclic graph)

- Tree is easiest for search, DCG hardest

- *Game graph*, *game tree* are other terms for state space of games



Examples of Types of State Spaces

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces



Same Go state
reached with different
move order

- Tic-Tac-Toe: DAG
 - Different move order can lead to same result
- Go without repetition rules: DCG (cycles exist, e.g. simple ko)
- Go with simple ko rules: still DCG (longer cycles still exist)
- Go with full repetition rules: DAG (details later)
- Chess, checkers, NoGo: discuss.

Complexity of a State Space

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Some measures of game complexity:
- Size of state space
- Branching factor (number of actions in state)
- More on these later
- Difficulty of game can depend on many other things
 - Is there a simple strategy?
 - A mathematical theory?
 - Many master games to learn from?
 - Good heuristics?
- Difficulty for humans \neq difficulty for computers

Game Result

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- What is the result when a game is over?
- Simplest: win/loss
 - Go with non-integer komi, NoGo, Nim, ...
- Many games: win/loss/draw
 - Chess, checkers, tic-tac-toe, Go with integer komi, five-in-a-row (gomoku), ...
- Point-scoring games: size of win matters
 - Score, value, evaluation, result:
terms with similar meaning

Types of Game Representation

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Board representation
- Move history - move list
- Formats:
 - Internal representation
 - File storage: move list + annotations, e.g. sgf file format
 - www.red-bean.com/sgf/
 - en.wikipedia.org/wiki/Smart_Game_Format
 - Inter-program communication: GTP - Go Text Protocol
 - www.lysator.liu.se/~gunnar/gtp/
 - Used to connect Go0 etc to GoGui and other tools

Review - Board Representation

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Many games are grid-based
- 2-d array
- 1-d array (often faster, standard)
- Bitmaps (sometimes fastest, depends on use case)
- Specialized, e.g. piece list if large board, few pieces.

Summary - Concepts so Far

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- State space search - formal model for decision-making
- Basic general concepts:
states, actions, state space
- In games:
game position, game state, move, game tree/game graph
- Terminal states, result, reward
- Representation of board and moves, internal and external (file storage, text communication)

Models of State Spaces

Models of State Spaces

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- How to choose a state space for a decision-making problem?
- More on Tree, DAG, DCG
- Estimating the size of state space
- Reachable vs unreachable
- Symmetries, equivalent states

How to Choose States and State Spaces

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- State needs all the *relevant information*
- Make it as simple as possible ...
- ... but not simpler

Example 1: Simplifying Rewards in States

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Example: a zero-sum game, both players collect rewards during the game
- Most direct representation:
list of rewards at each move for each player
- Example of state:
(`[0, 10, 0, 0, 10]`, `[20, 0, -5, 0, 0]`,
`rest-of-state`)
- Better: just keep sum of awards so far:
(`20, 15, rest-of-state`)
- Even better: only keep *difference* of rewards
(`5, rest-of-state`)

Example 1 Discussion

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Why is this representation better?
- All states with same reward difference are *equivalent*
 - The players do not care which of them they are in
- Great reduction in number of states
 - If one of the equivalent states is solved, then *all* are solved
 - Idea here: compress history into a single number, add that to state




Example 2: Simplifying State in Go

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces



- Consider the *simple ko* rule in Go
- Need to check if the position is the same as two moves ago
- Simple way: store the positions, compare each point on board
- Better way: work out the conditions when the simple ko rule applies
- We need to only store a single point.
 - See `ko_recapture` in `Go1`
- Idea here: compress history into information of a single relevant point, add that to state

More on State Spaces - Tree, DAG, DCG

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- What are the differences between tree, DAG, DCG?
- If we have a choice, which one to choose?
- We often have a choice!



State Space - Tree

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces



- Simplest state space model is tree
- Every action leads to new state
- Only one path from root to a node
- We simply ignore it, if different paths lead to the same situation
- Each copy is a separate state
- Example: Two sequences of Go moves
 - 1. Black B4, 2. White C3, 3. Black D1
 - 1. Black D1, 2. White C3, 3. Black B4
- Different sequences, different states in tree model. Duplication!

Pros and Cons of Tree

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

Advantages

- Simplest model
- Single path to each node
- No dependencies

Disadvantages

- Duplication, no re-use of information
- State space can be much larger than needed
- Search can become very inefficient, searches many copies of equivalent sub-trees

Size of State Space - Simple Model

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- How many states?
- Model:
 - Assume a *constant branching factor* b
 - Each *interior* node in tree has b children
 - Assume a *uniform depth* d
 - Each path from root is d actions long

State Space - Simple Model

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- How many nodes?
- 1 root node, $1 = b^0$ total nodes at depth 0
- b children of root, b^1 total nodes at depth 1
- Each child has b new children, total b^2 at depth 2
- ...
- Last level b^n nodes at depth n
- Total nodes $1 + b + ...b^n = (b^{n+1} - 1)/(b - 1)$
- For large b , this is close to b^n - last level dominates

(Bad) Example - Tic Tac Toe

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- 9 possible moves, $b = 9$
- depth at most 9, $d = 9$
- So about $9^9 \approx 387$ million nodes?
- This is a huge over-estimate, why?
- (Discuss)

A Better Model?

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- Games such as TicTacToe only add stones, never remove
- n choices for first move
- $n - 1$ for second move, ...
- Total $n \times (n - 1) \times \dots 1 = n!$ possible games
- TicTacToe: $9! = 362,880$
- Still too much. Discuss why.

(Very rough) Example - 7×7 Go

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- 49 moves at start, maybe 25 on average during a game
- Length of game - about 30 moves?
- Rough estimate $25^{30} \approx 10^{42}$
- Compare with $49! \approx 10^{63}$

Summary - Models of State Spaces

CMPUT 455

Lecture 4:
Formalizing
Decision-
Making

Formalizing
State Space
Search

Models of
State Spaces

- State space search - formal model for decision-making
- Basic general concepts: states, actions, state space
- In games: game position, game state, move, game tree/game graph
- Terminal states, result, reward
- Representation of board and moves, internal and external